

<> Code

🕒 Issues 3

🔗 Pull requests

▶ Actions

📁 Projects


🛡 Security

📈 Insights

🔑 main ▾

⋮

100-days-of-machine-learning / day59-classification-metrics / classification-metrics-binary.ipynb

 campusx-official Add files via upload

🕒 History

👤 1 contributor

1 lines (1 sloc) | 19.5 KB

⋮

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets mounted as /kaggle/working.
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session

/kaggle/input/heart-disease-uci/heart.csv
```

```
In [2]: df = pd.read_csv('/kaggle/input/heart-disease-uci/heart.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	ta
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	

```
In [4]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(df.iloc[:,0:-1],df.iloc[:,15:-1],
```

```
In [5]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

```
In [6]: clf1 = LogisticRegression()
clf2 = DecisionTreeClassifier()
```

```
In [7]: clf1.fit(X_train,y_train)
clf2.fit(X_train,y_train)
```

/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:765:

```
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
Out[7]: DecisionTreeClassifier()
```

```
In [8]: y_pred1 = clf1.predict(X_test)
        y_pred2 = clf2.predict(X_test)
```

```
In [9]: from sklearn.metrics import accuracy_score, confusion_matrix
        print("Accuracy of Logistic Regression", accuracy_score(y_test, y_pred1))
        print("Accuracy of Decision Trees", accuracy_score(y_test, y_pred2))
```

```
Accuracy of Logistic Regression 0.9016393442622951
Accuracy of Decision Trees 0.8360655737704918
```

```
In [10]: confusion_matrix(y_test, y_pred1)
```

```
Out[10]: array([[26,  6],
               [ 0, 29]])
```

```
In [11]: print("Logistic Regression Confusion Matrix\n")
        pd.DataFrame(confusion_matrix(y_test, y_pred1), columns=list(range(0, 2)))
```

```
Logistic Regression Confusion Matrix
```

```
Out[11]:
```

	0	1
0	26	6
1	0	29

```
In [12]: print("Decision Tree Confusion Matrix\n")
        pd.DataFrame(confusion_matrix(y_test, y_pred2), columns=list(range(0, 2)))
```

```
Decision Tree Confusion Matrix
```

```
Out[12]:
```

	0	1
0	24	8
1	2	27

```
In [13]: result = pd.DataFrame()
        result['Actual Label'] = y_test
        result['Logistic Regression Prediction'] = y_pred1
        result['Decision Tree Prediction'] = y_pred2
```

In [14]: `result.sample(10)`

Out[14]:

	Actual Label	Logistic Regression Prediction	Decision Tree Prediction
257	0	0	0
65	1	1	1
164	1	1	1
29	1	1	1
296	0	1	0
184	0	0	0
161	1	1	1
292	0	0	0
53	1	1	1
147	1	1	1

In [15]: `from sklearn.metrics import recall_score, precision_score, f1_score`

In [16]:

```

print("For Logistic regression Model")
print("-"*50)
cdf = pd.DataFrame(confusion_matrix(y_test,y_pred1),columns=list(range(0,2)))
print(cdf)
print("-"*50)
print("Precision - ",precision_score(y_test,y_pred1))
print("Recall - ",recall_score(y_test,y_pred1))
print("F1 score - ",f1_score(y_test,y_pred1))

```

For Logistic regression Model

```

-----
      0    1
0  26    6
1    0   29
-----

```

```

Precision -  0.8285714285714286
Recall -    1.0
F1 score -  0.90625

```

In [17]:

```

print("For DT Model")
print("-"*50)
cdf = pd.DataFrame(confusion_matrix(y_test,y_pred2),columns=list(range(0,2)))
print(cdf)
print("-"*50)
print("Precision - ",precision_score(y_test,y_pred2))
print("Recall - ",recall_score(y_test,y_pred2))
print("F1 score - ",f1_score(y_test,y_pred2))

```

For DT Model

```

-----
      0    1
0  24    8
1    6   26
-----

```

```
0  4  0
1  2  27
```

```
-----
Precision - 0.7714285714285715
Recall - 0.9310344827586207
F1 score - 0.8437500000000001
```

```
In [18]: precision_score(y_test,y_pred1,average=None)
```

```
Out[18]: array([1.          , 0.82857143])
```

```
In [19]: precision_score(y_test,y_pred2,average=None)
```

```
Out[19]: array([0.92307692, 0.77142857])
```

```
In [21]: recall_score(y_test,y_pred2,average=None)
```

```
Out[21]: array([0.75          , 0.93103448])
```