
 campusx-official Add files via upload History

 1 contributor

```
In [42]: import numpy as np # linear algebra
import pandas as pd # data processing
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [43]: df = pd.read_csv('Social_Network_Ads.csv')
```

```
In [44]: df=df.iloc[:,2:]
```

```
In [45]: df.sample(5)
```

```
Out[45]:
```

	Age	EstimatedSalary	Purchased
137	30	107000	1
251	37	52000	0
262	55	125000	1
144	34	25000	0
292	55	39000	1

## Train test split

```
In [46]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df.drop('Purchased', axis=1),
                                                    df['Purchased'],
                                                    test_size=0.3,
                                                    random_state=0)

X_train.shape, X_test.shape
```

```
Out[46]: ((280, 2), (120, 2))
```

## StandardScaler

```
In [47]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

# fit the scaler to the train set, it will learn the parameters
scaler.fit(X_train)

# transform train and test sets
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [48]: scaler.mean_
```

```
Out[48]: array([3.78642857e+01, 6.98071429e+04])
```

```
In [49]: X_train
```

```
Out[49]:
```

	Age	EstimatedSalary
--	-----	-----------------

<b>92</b>	26	15000
-----------	----	-------

<b>223</b>	60	102000
------------	----	--------

<b>234</b>	38	112000
------------	----	--------

<b>232</b>	40	107000
------------	----	--------

<b>377</b>	42	53000
------------	----	-------

...	...	...
-----	-----	-----

<b>323</b>	48	30000
------------	----	-------

<b>192</b>	29	43000
------------	----	-------

<b>117</b>	36	52000
------------	----	-------

<b>47</b>	27	54000
-----------	----	-------

<b>172</b>	26	118000
------------	----	--------

280 rows × 2 columns

```
In [52]: X_train_scaled
```

```
Out[52]:
```

	Age	EstimatedSalary
--	-----	-----------------

<b>0</b>	-1.163172	-1.584970
----------	-----------	-----------

<b>1</b>	2.170181	0.930987
----------	----------	----------

<b>2</b>	0.013305	1.220177
----------	----------	----------

<b>3</b>	0.209385	1.075582
----------	----------	----------

<b>4</b>	0.405465	-0.486047
----------	----------	-----------

...	...	...
-----	-----	-----

<b>275</b>	0.993704	-1.151185
------------	----------	-----------

<b>276</b>	-0.869053	-0.775237
------------	-----------	-----------

<b>277</b>	-0.182774	-0.514966
------------	-----------	-----------

<b>278</b>	-1.065133	-0.457127
------------	-----------	-----------

<b>279</b>	-1.163172	1.393691
------------	-----------	----------

280 rows × 2 columns

```
In [51]: X_train_scaled = pd.DataFrame(X_train_scaled, columns=X_train.columns)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X_test.columns)
```

```
In [9]: np.round(X_train.describe(), 1)
```

```
Out[9]:
```

	Age	EstimatedSalary
<b>count</b>	280.0	280.0
<b>mean</b>	37.9	69807.1
<b>std</b>	10.2	34641.2
<b>min</b>	18.0	15000.0
<b>25%</b>	30.0	43000.0
<b>50%</b>	37.0	70500.0
<b>75%</b>	46.0	88000.0
<b>max</b>	60.0	150000.0

```
In [10]: np.round(X_train_scaled.describe(), 1)
```

```
Out[10]:
```

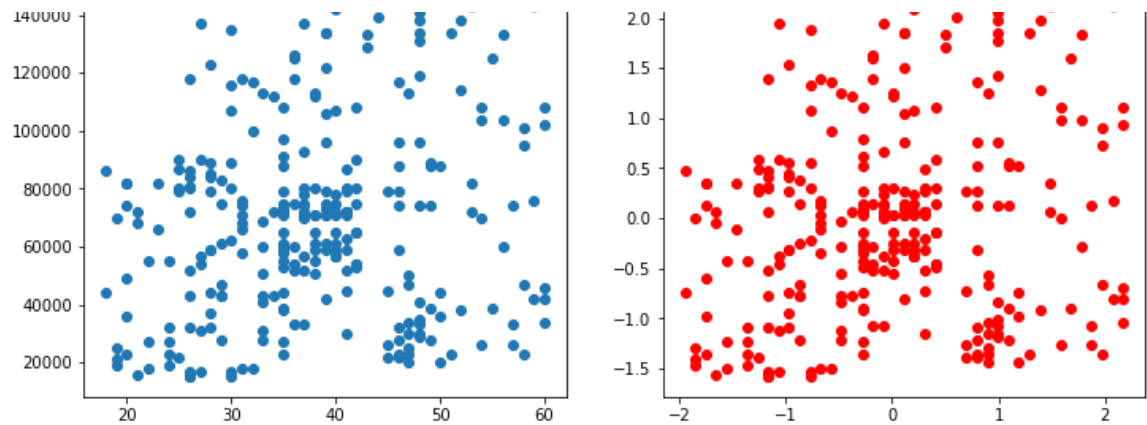
	Age	EstimatedSalary
<b>count</b>	280.0	280.0
<b>mean</b>	0.0	0.0
<b>std</b>	1.0	1.0
<b>min</b>	-1.9	-1.6
<b>25%</b>	-0.8	-0.8
<b>50%</b>	-0.1	0.0
<b>75%</b>	0.8	0.5
<b>max</b>	2.2	2.3

## Effect of Scaling

```
In [11]: fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 5))

ax1.scatter(X_train['Age'], X_train['EstimatedSalary'])
ax1.set_title("Before Scaling")
ax2.scatter(X_train_scaled['Age'], X_train_scaled['EstimatedSalary'], color='red')
ax2.set_title("After Scaling")
plt.show()
```



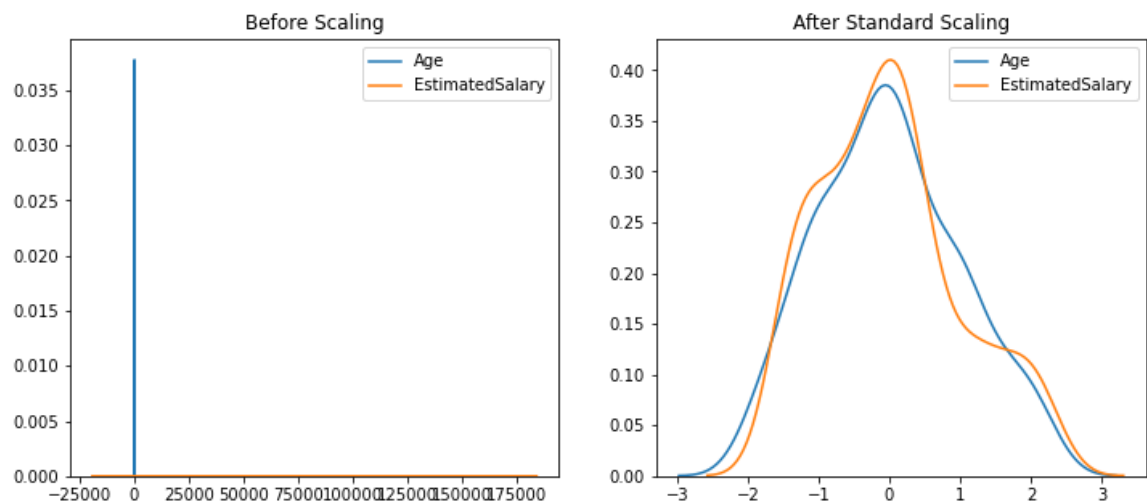


In [12]:

```
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 5))

# before scaling
ax1.set_title('Before Scaling')
sns.kdeplot(X_train['Age'], ax=ax1)
sns.kdeplot(X_train['EstimatedSalary'], ax=ax1)

# after scaling
ax2.set_title('After Standard Scaling')
sns.kdeplot(X_train_scaled['Age'], ax=ax2)
sns.kdeplot(X_train_scaled['EstimatedSalary'], ax=ax2)
plt.show()
```



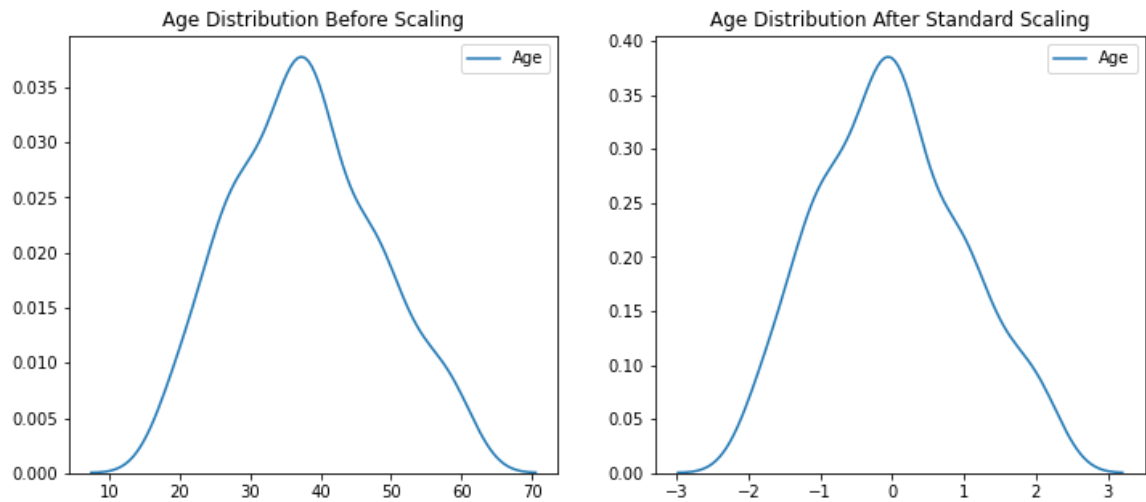
## Comparison of Distributions

In [13]:

```
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 5))

# before scaling
ax1.set_title('Age Distribution Before Scaling')
sns.kdeplot(X_train['Age'], ax=ax1)

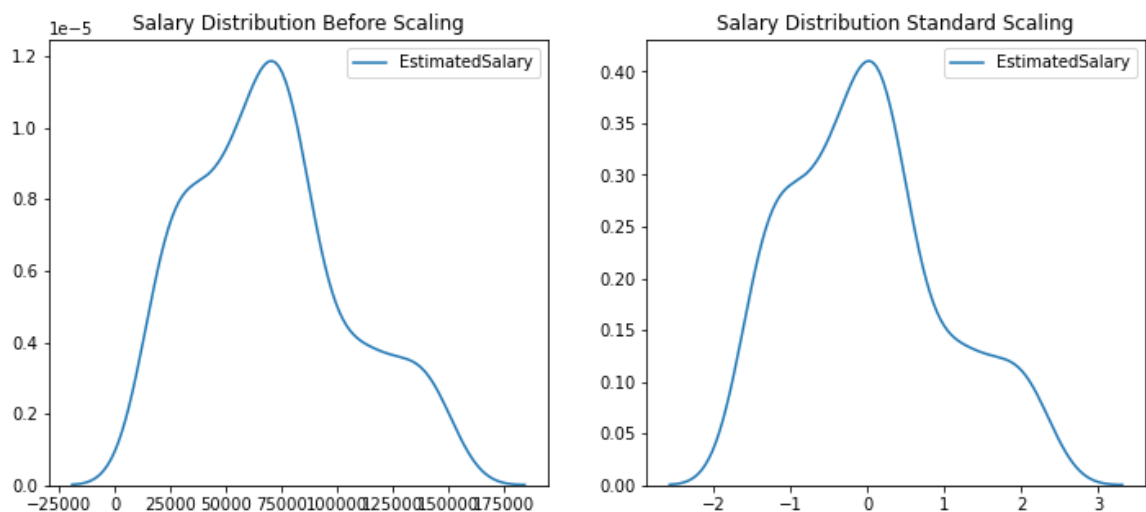
# after scaling
ax2.set_title('Age Distribution After Standard Scaling')
sns.kdeplot(X_train_scaled['Age'], ax=ax2)
plt.show()
```



```
In [14]: fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 5))

# before scaling
ax1.set_title('Salary Distribution Before Scaling')
sns.kdeplot(X_train['EstimatedSalary'], ax=ax1)

# after scaling
ax2.set_title('Salary Distribution Standard Scaling')
sns.kdeplot(X_train_scaled['EstimatedSalary'], ax=ax2)
plt.show()
```



## Why scaling is important?

```
In [15]: from sklearn.linear_model import LogisticRegression
```

```
In [17]: lr = LogisticRegression()
lr_scaled = LogisticRegression()
```

```
In [18]: lr.fit(X_train,y_train)
lr_scaled.fit(X_train_scaled,y_train)
```

```
lr_scaled = lr.predict(X_test_scaled, y_test_scaled)
```

Out[18]: LogisticRegression()

```
In [20]: y_pred = lr.predict(X_test)
y_pred_scaled = lr_scaled.predict(X_test_scaled)
```

```
In [21]: from sklearn.metrics import accuracy_score
```

```
In [22]: print("Actual", accuracy_score(y_test, y_pred))
print("Scaled", accuracy_score(y_test, y_pred_scaled))
```

Actual 0.6583333333333333  
Scaled 0.8666666666666667

```
In [23]: from sklearn.tree import DecisionTreeClassifier
```

```
In [24]: dt = DecisionTreeClassifier()
dt_scaled = DecisionTreeClassifier()
```

```
In [25]: dt.fit(X_train, y_train)
dt_scaled.fit(X_train_scaled, y_train)
```

Out[25]: DecisionTreeClassifier()

```
In [26]: y_pred = dt.predict(X_test)
y_pred_scaled = dt_scaled.predict(X_test_scaled)
```

```
In [27]: print("Actual", accuracy_score(y_test, y_pred))
print("Scaled", accuracy_score(y_test, y_pred_scaled))
```

Actual 0.875  
Scaled 0.875

```
In [29]: df.describe()
```

Out[29]:

	Age	EstimatedSalary	Purchased
<b>count</b>	400.000000	400.000000	400.000000
<b>mean</b>	37.655000	69742.500000	0.357500
<b>std</b>	10.482877	34096.960282	0.479864
<b>min</b>	18.000000	15000.000000	0.000000
<b>25%</b>	29.750000	43000.000000	0.000000
<b>50%</b>	37.000000	70000.000000	0.000000
<b>75%</b>	46.000000	88000.000000	1.000000

max 60.000000 150000.000000 1.000000

## Effect of Outlier

```
In [34]: df = df.append(pd.DataFrame({'Age':[5,90,95], 'EstimatedSalary':[1000,250000,350
```

```
In [32]: df
```


```
Out[32]:
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0
...	...	...	...
395	46	41000	1
396	51	23000	1
397	50	20000	1
398	36	33000	0
399	49	36000	1

400 rows × 3 columns

```
In [36]: plt.scatter(df['Age'], df['EstimatedSalary'])
```

```
Out[36]:
```



```
In [37]: from sklearn.model_selection import train_test_split
```



```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df.drop('Purchased', axis=1),
                                                    df['Purchased'],
                                                    test_size=0.3,
                                                    random_state=0)

X_train.shape, X_test.shape

```

Out[37]: ((282, 2), (121, 2))

```

In [38]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

# fit the scaler to the train set, it will learn the parameters
scaler.fit(X_train)

# transform train and test sets
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

```

In [40]: X_train_scaled = pd.DataFrame(X_train_scaled, columns=X_train.columns)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X_test.columns)


```

```

In [41]: fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 5))

ax1.scatter(X_train['Age'], X_train['EstimatedSalary'])
ax1.set_title("Before Scaling")
ax2.scatter(X_train_scaled['Age'], X_train_scaled['EstimatedSalary'], color='red')
ax2.set_title("After Scaling")
plt.show()

```



In [ ]:

